

Introduction to CFEngine 3

Riccardo Murri <riccardo.murri@uzh.ch>

What is CFEngine?

“CFEngine is a software solution that helps system administrators and other stakeholders in the IT organization become more agile and respond faster to business requirements while ensuring SLAs and regulatory compliance, through automation.” —

<http://cfengine.com/learn/what-is-cfengine/>

Huh? What the heck is CFEngine?

“Cfengine, or the configuration engine is an agent / software robot and a high level policy language [...] to administrate and configure large computer networks.” — CFEngine v2, <https://www.gnu.org/software/cfengine/>

“For many users, CFEngine is simply a configuration tool – i.e. software for deploying and patching systems according to a policy.” — CFEngine v3, <https://docs.cfengine.com/latest/guide-introduction.html>

Disclaimers

1. CFEngine works on GNU/Linux, *BSD, Solaris, and Windows. *My experience is limited to GNU/Linux only.*
2. CFEngine comes in a free/open-source version (“CFEngine community”) and a commercial version (“CFEngine Enterprise”), with more features and professional support. *My experience is limited to the “CFEngine community” only.*

Alternatives?

Puppet — a better CFEngine 2 ;-), written and extensible in Ruby.

Chef — a Ruby framework for managing systems.

SaltStack — A configuration management and remote execution framework. Conceptually very similar to Puppet, but written and extensible in Python.

Ansible — not quite an alternative: it's more geared to software deployment than configuration management. . .

(Theory)

So, how does it work?

*“The idea of CFEngine is to create a single file or set of configuration files which will describe the setup of every host on your network. [...] the **configuration of the host is checked against this model and, if necessary, any deviations are fixed.**”¹*

The host configuration (“model”) is built from atoms called **promises**.

¹<https://www.gnu.org/software/cfengine/cfdetails.html>

The agent workflow

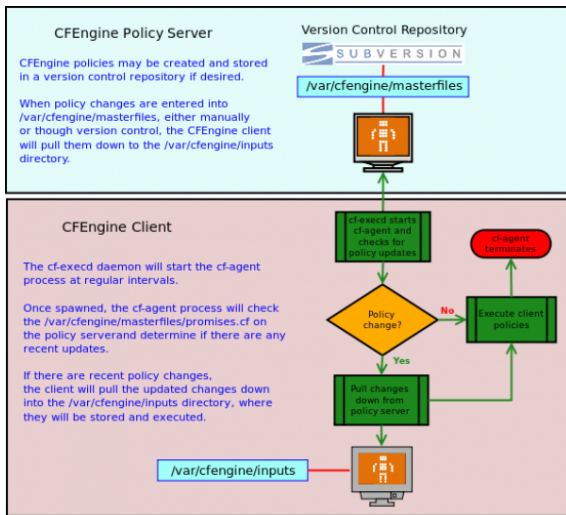


Image source: <http://cfengine.com/>

Promises

“A promise is the documentation or definition of an intention to act or behave in some manner. They are the rules which CFEngine clients are responsible for implementing.”²

```
files:
  "/tmp/email.[a-z0-9]+"
  file_select    => days_old("10"),
  delete         => tidy,
  comment        => "Delete temporary email files";
```

²<https://docs.cfengine.com/latest/guide-writing-and-serving-policy.html>

Promises

“A promise is the documentation or definition of an intention to act or behave in some manner. They are the rules which CFEngine clients are responsible for implementing.”²

```
files:  
  "/tmp/email.[a-z0-9]+"  
  file_select    => days_old("10"),  
  delete        => tidy,  
  comment       => "Delete temporary email files";
```

Target of the promise. A set of files in this case.

²<https://docs.cfengine.com/latest/guide-writing-and-serving-policy.html>

Promises

“A promise is the documentation or definition of an intention to act or behave in some manner. They are the rules which CFEngine clients are responsible for implementing.”²

```
files:
  "/tmp/email.[a-z0-9]+"
  file_select    => days_old("10"),
  delete         => tidy,
  comment        => "Delete temporary email files";
```

Additional criteria for building the promise target.

²<https://docs.cfengine.com/latest/guide-writing-and-serving-policy.html>

Promises

“A promise is the documentation or definition of an intention to act or behave in some manner. They are the rules which CFEngine clients are responsible for implementing.”²

```
files:
  "/tmp/email.[a-z0-9]+"
  file_select    => days_old("10"),
  delete        => tidy,
  comment       => "Delete temporary email files";
```

Action to take. In this case, delete all (“tidy”).

²<https://docs.cfengine.com/latest/guide-writing-and-serving-policy.html>

Promises

A promise has three possible outcomes:

state is...	▼state was...▼	
▼	as promised	not as such
as promised	kept	repaired
not as such	<i>(impossible!)</i>	not kept

Bundles, I

Promises must be grouped into **bundles**.

```
bundle agent test {  
  vars:  
    "secrets" slist => { "/etc/shadow",  
                          "/etc/sudoers" };  
  
  files:  
    "$(secrets) "  
    create => "false",  
    perms  => mog("0400", "root", "root");  
}
```

What CFEngine component should execute this bundle. More on this later.

CFEngine 3 components




Image source: <http://cfengine.com/>

(Practice)

Example: keep secret files only readable by root

```
bundle agent keep_secrets {  
  vars:  
    "secrets" slist => { "/etc/shadow",  
                          "/etc/sudoers" };  
  
  files:  
    "$(secrets)"  
    create => "false",  
    perms  => mog("0400", "root", "root");  
}
```

Define secrets as a *list of strings*. Also available: simple strings, and (fake) integers and floats.

Example: keep secret files only readable by root

```
bundle agent keep_secrets {  
  vars:  
    "secrets" slist => { "/etc/shadow",  
                          "/etc/sudoers" };  
  files:  
    "$(secrets)"  
    create => "false",  
    perms  => mog("0400", "root", "root");  
}
```

Since `secrets` is a *list*, this implicitly *loops over all its values*.

Example: keep secret files only readable by root

```
bundle agent keep_secrets {  
  vars:  
    "secrets" slist => { "/etc/shadow",  
                          "/etc/sudoers" };  
  
  files:  
    "$(secrets)"  
    create => "false",  
    perms  => mog("0400", "root", "root");  
}
```

`mog()` = mode-owner-group. It's a standard library *body*.

Bodies

“the definition of a promise can grow complicated. Complex promises are best understood by breaking them down into independent, re-usable components. The CFEngine reserved word `body` is used to encapsulate the details of complex promise attribute values. Bodies can optionally have parameters.”

```
body perms secret {  
    mode    => "0600";  
    owners => { "root" };  
    groups => { "root" };  
}
```

```
body perms  
    mog(mode, user, group)  
{  
    owners=>{"$(user)"};  
    groups=>{"$(group)"};  
    mode  =>"$(mode)";  
}
```

Reference:

<https://docs.cfengine.com/latest/guide-language-concepts-bodies.html>

Example: edit a file

```
bundle agent nss_myhostname {  
  files:  
    "/etc/nsswitch.conf"  
    edit_line => enable_nss_myhostname;  
}  
  
bundle edit_line enable_nss_myhostname {  
  replace_patterns:  
    "(?x) ^hosts:\s* (((?!myhostname)\w+\s+)+) $"  
    replace_with => value(  
      "hosts: $(match.1) myhostname";  
    )  
}
```

Editing is a complex operation, that requires a sequence of action. **Therefore, it has a special kind of bundle.**

Example: edit a file

```
bundle agent nss_myhostname {
  files:
    "/etc/nsswitch.conf"
    edit_line => enable_nss_myhostname;
}

bundle edit_line enable_nss_myhostname {
  replace_patterns:
    "(?x) ^hosts:\s* (((?!myhostname)\w+\s+)+) $"
    replace_with => value(
      "hosts: $(match.1) myhostname");
}
```

Full **PCRE** syntax is supported everywhere in CFEngine.

Example: edit a file

```
bundle agent nss_myhostname {
  files:
    "/etc/nsswitch.conf"
    edit_line => enable_nss_myhostname;
}

bundle edit_line enable_nss_myhostname {
  replace_patterns:
    "(?x) ^hosts:\s* (((?!myhostname)\w+\s+)+) $"
    replace_with => value(
      "hosts: $(match.1) myhostname");
}
```

CFEngine will not let you substitute a regexp with a string that matches the same regexp! Need to learn what “negative lookahead assertions” are...

Example: edit a file

```
bundle agent nss_myhostname {  
  files:  
    "/etc/nsswitch.conf"  
    edit_line => enable_nss_myhostname;  
}  
  
bundle edit_line enable_nss_myhostname {  
  replace_patterns:  
    "(?x) ^hosts:\s* (((?!myhostname)\w+\s+)+) $"  
    replace_with => value(  
      "hosts: $(match.1) myhostname");  
}
```

Also available: `insert_lines`, `delete_lines`, `field_edits`. Which can all be combined together.

Contexts / Classes

Contexts³ are just boolean **constants**. The context “any” is CFEngine’s alias for “true”.

CFEngine has a concise syntax for “if” clauses using contexts:

```
vars:
  os_centos|os_rhel::
    "rules"          string => "/etc/sysconfig/iptables";
  os_debian|os_ubuntu::
    "rules"          string => "/etc/iptables/rules.v4";
```

³Previously called “Classes”

Contexts / Classes

Contexts³ are just boolean **constants**. The context “any” is CFEngine’s alias for “true”.

CFEngine has a concise syntax for “if” clauses using contexts:

```
vars:  
  os_centos|os_rhel::  
    "rules"          string => "/etc/sysconfig/iptables";  
  os_debian|os_ubuntu::  
    "rules"          string => "/etc/iptables/rules.v4";
```

This part will only be executed when the context `os_centos` or the context `os_rhel` have the “true” value.

³Previously called “Classes”

Example: deploy template file, I

```
bundle agent iptables {  
  vars:  
    os_centos|os_rhel::  
      "rules"      string => "/etc/sysconfig/iptables";  
    os_debian|os_ubuntu::  
      "rules"      string => "/etc/iptables/rules.v4";  
    any::  
      "www_ports" ilist => { "80", "443" };  
  files:  
    "$(rules)"  
    edit_template => "iptables.tpl",  
    edit_defaults => empty,  
    perms         => secret,  
    create        => "true";  
}
```

What template file to use.

Example: deploy template file, I

```
bundle agent iptables {
  vars:
    os_centos|os_rhel::
      "rules"      string => "/etc/sysconfig/iptables";
    os_debian|os_ubuntu::
      "rules"      string => "/etc/iptables/rules.v4";
    any::
      "www_ports" ilist => { "80", "443" };
  files:
    "$(rules)"
    edit_template  => "iptables.tpl",
    edit_defaults  => empty,
    perms          => secret,
    create         => "true";
}
```

Initial content of the editing buffer. In other words:
what to do if the destination file already exists, scratch
or append?

Example: deploy template file, I

```
bundle agent iptables {  
  vars:  
    os_centos|os_rhel::  
      "rules"      string => "/etc/sysconfig/iptables";  
    os_debian|os_ubuntu::  
      "rules"      string => "/etc/iptables/rules.v4";  
  any::  
    "www_ports" ilist => { "80", "443" };  
  files:  
    "$(rules)"  
    edit_template  => "iptables.tmpl",  
    edit_defaults  => empty,  
    perms          => secret,  
    create         => "true";  
}
```

Needed to un-do the limiting effect of context
`os_debian` up above.

Example: deploy template file, II

The actual contents of the template file:

```
# iptables.tmpl
[% CFEngine www:: %]
-A INPUT -p tcp -dport $(www_ports) -j ACCEPT
[% CFEngine am_policy_hub:: %]
-A INPUT -p tcp -dport 5308 -j ACCEPT
```

This is the way contexts are applied to templates. All subsequent lines are inserted if and only if context `www` is defined.

Example: deploy template file, II

The actual contents of the template file:

```
# iptables.tmpl
[% CFEngine www:: %]
-A INPUT -p tcp -dport $(www_ports) -j ACCEPT
[% CFEngine am_policy_hub:: %]
-A INPUT -p tcp -dport 5308 -j ACCEPT
```

Since `www_ports` is a list, **this implicitly loops over all values.**

Example: deploy template file, II

The actual contents of the template file:

```
# iptables.tmpl
[% CFEngine www:: %]
-A INPUT -p tcp -dport $(www_ports) -j ACCEPT
[% CFEngine am_policy_hub:: %]
-A INPUT -p tcp -dport 5308 -j ACCEPT
```

CFEngine defines this context on the policy servers.

Contexts / Classes, II

Additionally, contexts can be defined with specific promises, which evaluate a boolean expression:

```
classes:  
  "os_debian"  
  expression => "debian.!ubuntu";  
  
  "os_debian6"  
  expression => "os_debian&(debian_6|debian_squeeze)";
```

Class os_debian will be defined if class debian is defined and class ubuntu is not.

Contexts / Classes, II

Additionally, contexts can be defined with specific promises, which evaluate a boolean expression:

```
classes:  
  "os_debian"  
  expression => "debian.!ubuntu";  
  
  "os_debian6"  
  expression => "os_debian&(debian_6|debian_squeeze)";
```

The boolean *AND* operator can be spelled as dot “.” or with the more familiar “&”.

Contexts / Classes, II

Additionally, contexts can be defined with specific promises, which evaluate a boolean expression:

```
classes:  
  "os_debian"  
  expression => "debian. !ubuntu";  
  
  "os_debian6"  
  expression => "os_debian & (debian_6 | debian_squeeze)";
```

One of the quirks of CFEngine: **context debian is defined on Ubuntu systems as well.**

Example: install an updated package

```
bundle agent cve_2014_0160 {
  vars:
    "to_update" slist => {"openssl", "libssl1.0.0"};
  ubuntu_12_04::
    "ok_version" string => "1.0.1-4ubuntu5.12";
  debian_wheezy::
    "ok_version" string => "1.0.1e-2+deb7u5";

  packages:
    "$(to_update) "
    package_policy   => "update",
    package_version  => "$(ok_version)",
    package_select   => "<=",
    package_method   => generic;
}
```

Action to perform. Also available: install, remove, etc.

Example: install an updated package

```
bundle agent cve_2014_0160 {  
  vars:  
    "to_update" slist => {"openssl", "libssl1.0.0"};  
    ubuntu_12_04::  
      "ok_version" string => "1.0.1-4ubuntu5.12";  
    debian_wheezy::  
      "ok_version" string => "1.0.1e-2+deb7u5";  
  
  packages:  
    "$(to_update) "  
    package_policy   => "update",  
    package_version  => "$(ok_version)",  
    package_select   => "<=",  
    package_method   => generic;  
}
```

When to perform the action. Read: if the installed version is \leq of the promised version.

Example: install an updated package

```
bundle agent cve_2014_0160 {
  vars:
    "to_update" slist => {"openssl", "libssl1.0.0"};
  ubuntu_12_04::
    "ok_version" string => "1.0.1-4ubuntu5.12";
  debian_wheezy::
    "ok_version" string => "1.0.1e-2+deb7u5";

  packages:
    "$(to_update) "
    package_policy   => "update",
    package_version  => "$(ok_version)",
    package_select   => "<=",
    package_method   => generic;
}
```

Let CFEngine figure out what package manager to use.
Or you can be explicit and say: apt, yum, rpm, etc.

Contexts / Classes, III

Contexts can be set depending on the outcome of single promises:

```
files:
  "/etc/locale.gen"
  create      => "false",
  edit_line  => uncomment_lines_matching(
                        "(de|en|fr|it)_.+", "#"),
  classes    => if_repaired("run_localegen");
```

Context `run_locale_gen` will be set if this promise is “repaired”. One should use it to conditionally run the `locale-gen` command:

```
commands:
  os_debian.run_localegen::
    "/usr/sbin/locale-gen";
```

Example: extend CFEngine via modules

CFEngine can read back the output of any command and set/unset contexts or variables depending on it.

```
bundle local_ctx {  
  commands:  
    "cat locals.txt"  
    module => "true";  
}
```

```
# /var/cfengine/locals.txt  
+pizza_box  
=location[rack]=10  
=location[height]=2
```

Instructs CFEngine to read and parse the command output.

Example: extend CFEngine via modules

CFEngine can read back the output of any command and set/unset contexts or variables depending on it.

```
bundle local_ctx {  
  commands:  
    "cat locals.txt"  
    module => "true";  
}  
# /var/cfengine/locals.txt  
+pizza_box  
=location[rack]=10  
=location[height]=2
```

Sets context `pizza_box`

Example: extend CFEngine via modules

CFEngine can read back the output of any command and set/unset contexts or variables depending on it.

```
bundle local_ctx {  
  commands:  
    "cat locals.txt"  
    module => "true";  
}
```

```
# /var/cfengine/locals.txt  
+pizza_box  
=location[rack]=10  
=location[height]=2
```

Sets variables `location[rack]` and `location[height]`

Example: extend CFEngine via modules

CFEngine can read back the output of any command and set/unset contexts or variables depending on it.

```
bundle local_ctx {  
  commands:  
    "cat locals.txt"  
    module => "true";  
}  
# /var/cfengine/locals.txt  
+pizza_box  
=location[rack]=10  
=location[height]=2
```

Any line not starting with +, -, or = is ignored.

Normal ordering

“Within a bundle, the promise types are executed in a round-robin fashion according to so-called normal ordering [...]. The actual sequence continues for up to three iterations of the following, converging towards a final state: meta, vars, defaults, classes, users, files, packages, guest_environments, methods, processes, services, commands, storage, databases, reports”

Promises in a bundle are **not** executed in the order they are written; rather CFEngine groups them by type and executes types in a fixed order.

Reference: <https://docs.cfengine.com/latest/guide-language-concepts-normal-ordering.html>

Normal ordering

“Within a bundle, the promise types are executed in a round-robin fashion according to so-called normal ordering [...]. The actual sequence continues for up to three iterations of the following, converging towards a final state: meta, vars, defaults, classes, users, files, packages, guest_environments, methods, processes, services, commands, storage, databases, reports”

Note this! In other words, each promise may be re-evaluated up to three times during each bundle run.

Reference: <https://docs.cfengine.com/latest/guide-language-concepts-normal-ordering.html>

Example: send email

```
bundle agent send_email(to, subj, body) {  
  vars:  
    "tmpfile"  
    string => execresult("/bin/mktemp", "noshell");  
  methods:  
    "Write body contents into temp file"  
    usebundle => append_to_file("${tmpfile}",  
                                "${body}");  
  commands:  
    "/usr/bin/mail $(to) -s $(subj) <${tmpfile}";  
    "/bin/rm -f '${tmpfile}'";  
}
```

Example: send email

```
bundle agent send_email(to, subj, body) {  
  vars:  
    "tmpfile"  
    string => execresult("/bin/mktemp", "noshell");  
  methods:  
    "Write body contents into temp file"  
    usebundle => append_to_file("${tmpfile}",  
                                "${body}");  
  commands:  
    "/usr/bin/mail $(to) -s $(subj) <${tmpfile}";  
    "/bin/rm -f '${tmpfile}'";  
}
```

This results in a new file name at each pass. Hence, the bundle never “converges” and three emails are sent!

What other promise types are there?

- meta** Information about the bundle, mainly useful for documentation purposes.
- vars** *Definition of variables*
- defaults** “Default” values for bundle parameters
- classes** Definition of additional contexts
- users** Create or delete users in `/etc/passwd`
- files** *Create, edit, or audit files*
- packages** *Install or remove software packages*
- guest_environments** Control VMs using `libvirt`
- methods** Invoke other bundles
- processes** Stop or signal running processes.
- services** Start or stop system services.
- commands** *Run arbitrary commands.*
- storage** Mount NFS filesystems.
- databases** Create, alter, or manage DBs and tables on SQL, LDAP, or MS Registry
- reports** Print lines to the log file.

(The good, the bad, and the ugly) *.reverse()*

The ugly

It's (slowly) becoming a programming language, but the syntax is verbose and awkward, and is not getting better with time and releases:

```
bundle agent sysctl_data {  
  vars:  
    "parms_vars[net.ipv4.tcp_tw_reuse]" string => "1";  
    "parms_test_file" string => "/etc/sysctl";  
    "parms_debug" string => "on";  
    "parms_mgmt_policy" string => "ensure_present";  
  methods:  
    "test" usebundle => sysctl("sysctl_data.parms_");  
}
```

The bad

After more than 6 years, there are still serious bugs in the core functionality.

```
$ cat test.cf
bundle agent test {
  vars:
    "variables" slist => variablesmatching(".*");
  reports:
    linux::
      "$ (variables) = $ ( $ (variables) ) ";
}
```

```
$ time cf-agent -f ./test.cf -b test
[...]  
real 43m5.345s
```

That, and “normal ordering” which is the root of all evil.

The good

Strives to provide *idempotent primitives*, upon which to build your own systems administration DSL.

Can react to changing conditions: facts are not gathered at the beginning of the run.

Can *really* be extended using any language.

Runs 3'200 (and counting) promises on the top of every hour on $O(100)$ hosts in the S³IT server room, and *keeps all of them*. (Really, it can deploy and configure the entire OpenStack software suite at a whim.)

Appendix

Further reading

Learning CFEngine 3 The one and only beginners' book.

<https://docs.cfengine.com/docs/3.6/reference-promise-types.html>

The Reference Manual, you need this every time. . .


<https://docs.cfengine.com/latest/guide-introduction.html> What
the various components are, and how they interact.

<https://docs.cfengine.com/latest/guide-language-concepts.html>
Overview of the CFEngine language syntax.

<https://docs.cfengine.com/latest/guide-writing-and-serving-policy-promise>
Overview of the promise types and best practices for
structuring a CFEngine code base.

<https://docs.cfengine.com/latest/guide-design-center-configure-sketches->
Design Center / Sketches: how to re-use CFEngine
code packaged and shared by others.

CFEngine Puppet Salt Stack Ansible Chef



Google trends: 2004–today (▲), 2009–today (▼)

